
1 True or False? (40 points.) This question occupies 3 pages.

2 points for each correct answer.

1. On 16-bit machines, the largest memory size is 256 bytes.

ANSWER: False

2. In both 32-bit machines and 64-bit machines, the char pointers have the same length (in number of bytes).

ANSWER: False

3. The two's complement representation can encode the same number of negative integers and non-negative integers.

ANSWER: True

4. After running the following code snippet on a big-endian machine, the short variable p will hold the value "0000".

```
1      char x = 0x82;  
2      int y = (int)x;  
3      short p = *(&y);
```

ANSWER: False. It will be "ffff".

5. A program compiled on one machine can still run on another machine that supports a different instruction set architecture.

ANSWER: False

6. Executing the following two instructions sequentially (without anything happening in between) will not change the content in memory.

```
1      movq (%rax), %rdx  
2      movq %rdx, (%rax)
```

ANSWER: True

7. If after the following instructions **%rcx** holds the value of variable x, then **%rax** holds the value of &x.

```
1      movq (%rax), %rdx
2      movq (%rdx), %rcx
```

ANSWER: False. rdx holds the value &x.

8. To implement “for” loops, one must use both conditional jumps and unconditional jumps.

ANSWER: False. “Guided do” can be implemented with conditional jumps only.

9. Conditional jump instructions (jl, jne, etc.) use values in the general-purpose registers, in addition to the condition flags.

ANSWER: False

10. By convention, right before a callee function returns (before executing ret instruction), the stack pointer must always point to the top of the caller frame.

ANSWER: True

11. By convention, when a function returns, the callee function must explicitly let the caller function know which register holds the returned value.

ANSWER: False

12. By convention, if a caller function saves the value “5” in %rdi (a caller-saved register) right before entering a callee function, it can safely assume that when the callee returns, %rdi still holds the value “5”.

ANSWER: False

13. If a C struct and a C union have the same set of member fields, the size of the struct (in bytes) must be greater than or equal to the size of the union.

ANSWER: True

14. On a 64-bit machine, if we define the following struct, then after the struct is initialized in memory, the address of the short field is exactly 8 bytes away from the address of the long field.

```
1      struct S {
2          long a;
3          short b;
4      } *p
```

ANSWER: True. The alignment will ensure a starts at an address of multiple of 8, so b is automatically aligned without need of padding.

15. Some programs exhibit locality, and some don't.

ANSWER: False

16. Comparing a 6-way set associative cache to a 4-way set associative cache with the same capacity and block size, the 6-way cache generally results in less conflict misses, but it is more difficult to implement.

ANSWER: True

17. Miss rate + hit rate is always 100% with respect to a single cache and a single program.

ANSWER: True

18. nuX86_64 is exactly the same as x86_64, except that it also defines a software-managed cache and several new instructions to manage this cache. This new cache is a microarchitectural feature, and not an architectural (ISA) feature, like all other caches.

ANSWER: False

19. Consider a program that accesses every 4-byte element of an integer array in order, and then it accesses every element of the same array in order again for a second time. The machine implements a fully-associative cache with a block size of 8 bytes. The miss rate of the program on this machine is 25%.

ANSWER: False

20. In general, for caches, the write back policy generates less traffic to memory than the write through policy.

ANSWER: True

2 (8 points.)

2 points for each correct answer. If the answer is “No”, but the counter example is not correct, it will get only 1 point.

Assume that the length of `int` and `unsigned int` is 32 bits, signed integers use two’s complement, and right/left shifts on signed integers are arithmetic. The following code generates arbitrary values for `int x` and `int y`, then converts them to `unsigned int`.

```
1 int x = random();      /* Create some arbitrary values */
2 int y = random();
3 unsigned int ux = x;    /* Convert to unsigned */
4 unsigned int uy = y;
```

Each part below is either a C expression of equality (“==”) or a logical implication between two expressions ($A \implies B$). A non-zero value is true in C; zero means false.

Answer whether the expression or implication is *always* true for all possible `x` and `y` (type “Yes”), or not (type “No”). If “No”, give specific values for the variables involved to make it false. Specific values must be in hex (starting with “0x”), except that you may also use `TMin32` and `TMax32` for the smallest and largest values (respectively) that can be stored in an `int`.

A. $(x > 0) \ \&\& \ (y < 0) \implies (x-y > 0)$

No: Any `x-y` that overflows is a counter example.

B. $x > 11 \implies (x >> 2) == (\text{int})(ux >> 2)$

Yes

C. $x+0u \geq 0u$

Yes, because `x` will be treated as an unsigned `int`

D. $(\text{int})(ux + 5*y) == x + 5*y$

Yes (integer arithmetic is implemented same for signed and unsigned)

3 (16 points)

2 point for each correct answer.

Consider this pseudocode compiled on a machine for which the sizes (in bytes) of `char`, `short`, and `int` are 1, 2, and 4, respectively.

```
1  int main()
2  {
3      char c[4] = {0x12, 0x34, 0x56, 0x78};
4      char c1 = *((char *)(&c[2]));
5      short s1 = *((short *)(&c[0]));
6      short s2 = *((short *)(&c[1]));
7      int i1 = *((int *)(&c[0]));
8  }
```

A. What will be the following values on a big-endian machine (heximal format, starting with 0x)?

c1: _____ ANSWER: "0x56"

s1: _____ ANSWER: "0x1234"

s2: _____ ANSWER: "0x3456"

i1: _____ ANSWER: "0x12345678"

B. What will be the following values on a little-endian machine (heximal format, starting with 0x)?

c1: _____ ANSWER: "0x56"

s1: _____ ANSWER: "0x3412"

s2: _____ ANSWER: "0x5634"

i1: _____ ANSWER: "0x78563412"

4 (8 points.)

For parts A-C, 1 point for each correct answer. For part D, 2 points for the correct answer

The following values are stored at the indicated memory addresses and registers:

Address	Value	Register	Value
0x210	0xAB	%rax	0x210
0x218	0xBC	%rcx	0x220
0x220	0x09	%rdx	0x1
0x228	0x42	%rsi	0x2
0x230	0x54	%rdi	0x4

Assume that the values in the “Value” column occupy eight bytes of storage. For each part (A-D) of this question, answer the value stored in %r8, in heximal format (starting with 0x):

Part A:

```
1 movq (%rax), %r8
2 addq 0x6(%rax, %rsi), %r8
```

Value stored in %r8 after the 1st line is _____ ANSWER: "0xAB"

Value stored in %r8 after the 2nd line is _____ ANSWER: "0x167"

Part B:

```
1 movq 0x8(%rax), %r8
2 movq 0x6(%rcx, %rdx, 2), %r9
3 subq %r9, %r8
```

Value stored in %r8 after the 1st line is _____ ANSWER: "0xBC".

Value stored in %r8 after the 3rd line is _____ ANSWER: "0x7A".

Part C:

```
1 movq 0x228, %r8
2 leaq (%r8, %r8, 1), %r8
3 shrq %rdx, %r8
4 addq $0x228, %r8
```

Value stored in %r8 after the 2nd line is _____ ANSWER: "0x84"

Value stored in %r8 after the 4th line is _____ ANSWER: "0x26A"

Part D:

```
1 movq 0x230, %r8
2 movq $0x230, %r9
3 notq %r8
4 notq %r9
5 xorq %r9, %r8
```

Value stored in %r8 after the 5th line is _____ ANSWER: "0x264"

5 (14 points) This question occupies 2 pages.

5.1 The assembly code below was compiled from a C function whose prototype is given below.

```
long funcP(long r, s, t);
```

```
1 _funcQ: #long funcQ (long x, long y)
2     addq    $2, %rdi
3     imulq   $3, %rdi
4     movq    %rsi, %rax
5     addq    %rdi, %rax
6     ret
7
8 _funcP: #long funcP (long r, long s, long t)
9     movq    %rdi, %r8
10    subq    %rdx, %r8
11
12    testq    %r8, %r8
13    jle     foo
14
15    callq    _funcQ
16    jmp     bar
17
18 foo:
19    movq    %rsi, %rax
20    movq    %rdx, %rsi
21    movq    %rax, %rdx
22    callq    _funcQ
23
24 bar:
25    ret
```

Consider the assembly code above, assume that for parts A-D, the inputs are $r=2$, $s=3$, $t=5$.

A. How many times is funcQ called?

FuncQ is called _____ times. ANSWER: 1 time. 2 points for correct answer.

B. What are argument values, x and y, to funcQ when funcQ is called for the first time?

x = _____, y = _____ ANSWER: x = 2, y = 5. 2 points for each.

C. What is the return value of funcQ when it is called for the first time?

Return value of funcQ is _____ ANSWER: Answer is 17. 2 points for correct answer.

D. What is the return value of funcP?

Return value of funcP is: _____ ANSWER: Answer is 17. 3 points for correct answer.

5.2 The assembly code below was compiled from a C function whose prototype is given below.

```
int funcP(long[][] x);
```

```
1 _funcP: #int funcP (long[][] x)
2     movq $0, %rax
3     movq $0, %r8
4
5 .loop:
6     movq %r8, %rdx
7     imulq $0x10, %rdx
8     movq (%rdi, %rdx), %r9
9     movq 0x8(%rdi, %rdx), %r10
10
11     comq %r10, %r9
12     jle foo
13     addq %r9, %rax
14     jmp bar
15
16 .foo:
17     addq %r10, %rax
18
19 .bar:
20     incq %r8
21
22     comq $2, %r8
23     jle loop
24     ret
```

E. Consider the assembly code above, what will be the returned value, if the input is $\text{long}[3][2]x = \{\{1, 6\}, \{2, 3\}, \{5, 4\}\}$?

Return value is: _____ ANSWER: The return value is 14. 3 points for the correct answer

6 (14 Points)

Suppose we have a system with the following properties:

- The memory is byte addressable.
- Memory accesses are to **1-byte words** (not to 4-byte words).
- Addresses are 12 bits wide.
- The cache has four sets ($S = 4$). Each set consists of two lines (two-way set associative, $E = 2$). Each line holds four bytes of data ($B = 4$).
- LRU replacement is used for the cache.

The contents of the cache are as follows, with all addresses, tags, and values given in hex:

Set index	Valid	Tag	Byte 0	Byte 1	Byte 2	Byte 3
0	1	00	40	41	42	43
	1	83	FE	97	CC	D0
1	1	00	44	45	46	47
	0	83	—	—	—	—
2	1	00	48	49	4A	4B
	0	40	—	—	—	—
3	1	FF	9A	C0	03	FF
	0	54	—	—	—	—

Part A. For this configuration, which address bits are used for the tag, set index, and block offset? Indicate your answer by writing a string with “t” to indicate tag bits, “s” to indicate the index bits, and “b” to indicate block offset bits. For example, tttssssbbbbb.

Answer: _____

ANSWER: tttttttssbb; 2 points for correct answer

Part B. For each of the following memory reads which are **carried out in sequence** as listed, indicate two things:

1. Will it be a cache hit or miss?
2. If it is a hit, give the value of the corresponding memory location (in hex, starting with 0x). If it is a miss, write “unknown” in the “Value” column.

In total, there are 12 blanks to fill. The cache and memory are affected by nothing except the operations listed.

Address	Hit/Miss	Value
0x9	hit	0x49
0x54E	miss	unknown
0x836	miss	unknown
0xABD	miss	unknown
0x831	hit	0x97
0xFFF	miss	unknown or “0xFF”, since the value remains the same even after cache entry is evicted.

ANSWER: Answers are in the table. 1 point for each blank (12 blanks in total).

DO NOT DISTRIBUTE